CS5335 Final Project

By: Amit Mulay

Charles Denhart

Kelly King

Pablo Gomez

Original Goals: Maze Prison Escape

Implement a Maze Prison Escape Game.

Similar concept to PAC-MAN: main agent attempting to achieve some goals and a secondary agent acting as an adversary.

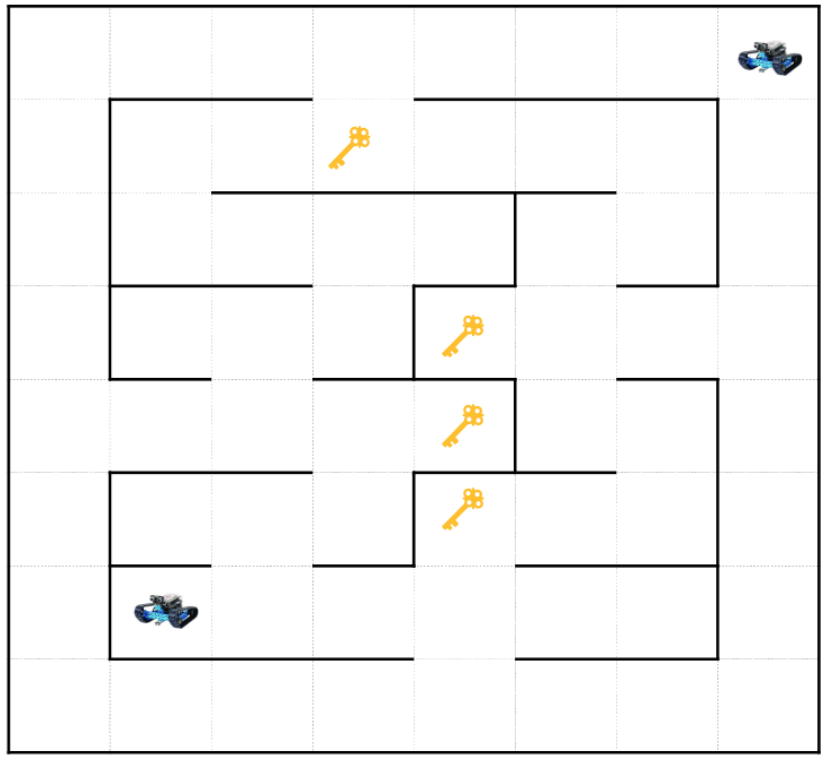

- Prisoner Goals/Behavior:
  - Escape Maze
  - Collect all the keys
  - React to lights and sound to avoid the patrol robot
- Patrol Goals/Behavior:
  - Patrol perimeter of the maze
  - Drive into the maze when it found an opening
- Maze configuration:
  - Beginning and End
  - At least 2 paths to each key

Considerations

- How to indicate the prisoner has escaped?
  - Map of the maze: robot stops when in escape position
- How will the Prisoner move around the maze?
  - Wall following would be the simplest solution to implement, but not conducive to avoiding the patrol's path.
- Method for collecting the keys?
  - Using the line follower sensor

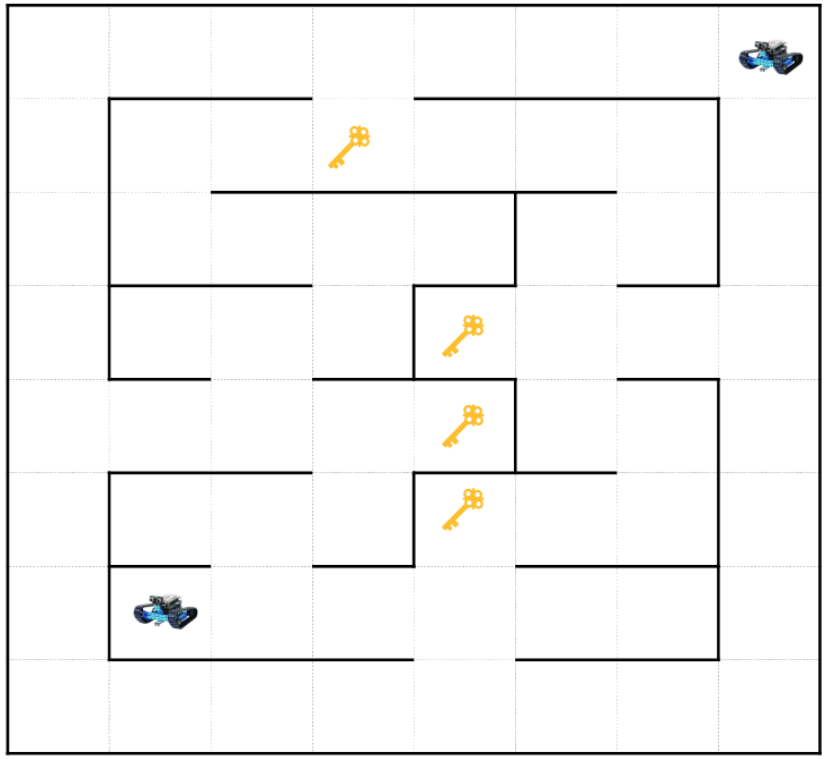

Considerations cont.

- How does the prisoner react to the patrol?
  - Patrol uses led-ring and buzzer to produce light and sound that the prisoner can detect
- How will the Patrol move?
  - Using the line follower sensor
- How will the patrol know to enter the maze?
  - Option A: Map the maze - turn and reverse at certain locations
  - Option B: Add an extra sensor - detect openings on the inside wall

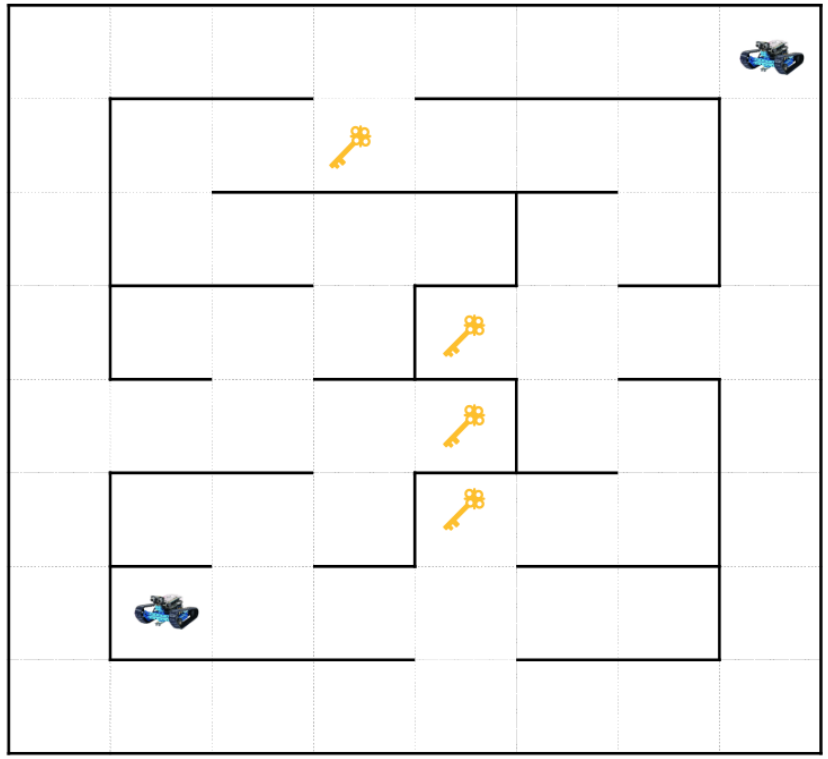

Original Goals: Maze Prison Escape

Implement a Maze Prison Escape Game.

Similar concept to PAC-MAN: main agent attempting to achieve some goals and a secondary agent acting as an adversary.

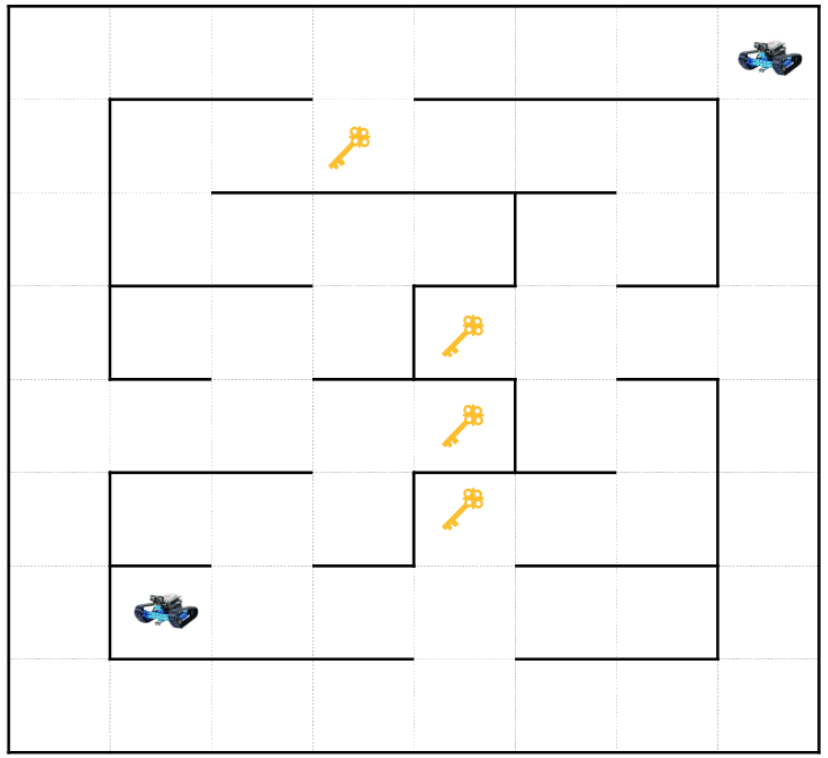

- Prisoner Goals/Behavior:
  - Escape Maze
  - Collect all the keys
  - React to lights and sound to avoid the patrol robot
- Patrol Goals/Behavior:
  - Patrol perimeter of the maze
  - Drive into the maze when it found an opening
- Maze configuration:
  - Beginning and End
  - At least 2 paths to each key

Revised Goals: Maze Prison Escape

Implement a Maze Prison Escape Game.

Similar concept to PAC-MAN: main agent attempting to achieve some goals and a secondary agent acting as an adversary.

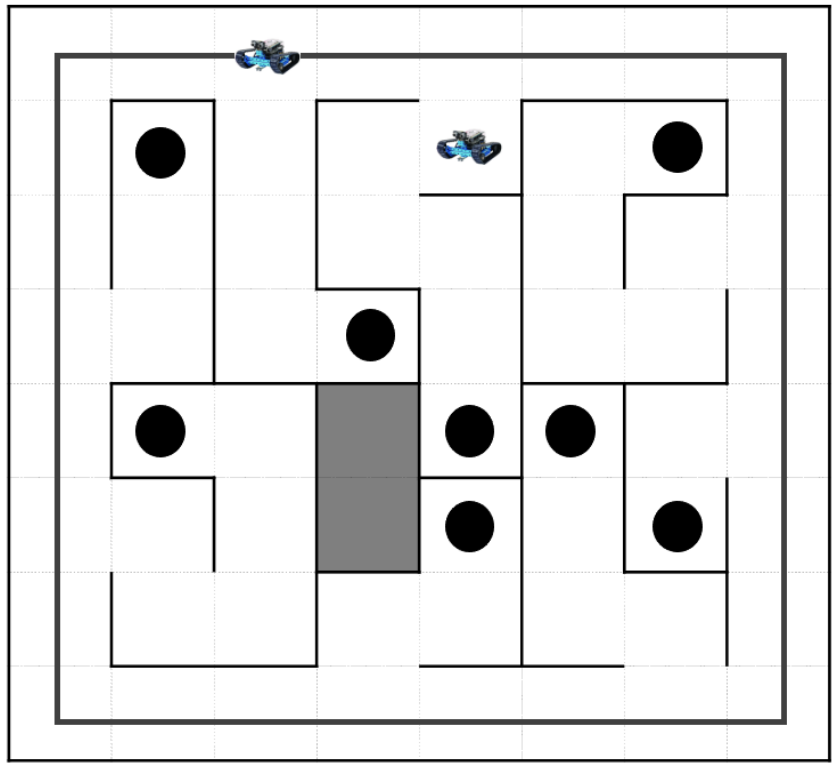

- Prisoner Goals and Behavior:
  - ~~Escape Maze~~
  - Collect all the keys
  - ~~React to lights and sound to avoid the patrol robot~~
- Patrol Goals:
  - Patrol perimeter of the maze
  - ~~Drive into the maze when it found an opening~~
  - Sleep and chase for randomized time intervals (10-20sec)
- Maze configuration:
  - ~~Beginning and End~~
  - ~~Keys are reachable without encountering Patrol~~
  - Keys are in "rooms" that require crossing the Patrols Path

Implementation

Prisoner

- [ ] Wall Following
- [ ] Reinforced Learning

Patrol

- [ ] Line Following
- [ ] Randomized Behavior

Academic Paper Technique:
Reinforcement Learning (RL)

- Q-Learning: $$Q^*(s,a) = E\left[R(s,a) + \gamma \max_{a'} Q^*(s',a')\right]$$
- Sparse reward function (vs. dense reward function)
- Uses original function approximation algorithm (HEDGER)
- Two tasks:
    1. Corridor following
    2. Obstacle avoidance with a target

Smart, W. D., & Kaelbling, L. P. (n.d.). Effective Reinforcement Learning for Mobile Robots. Retrieved from http://www.sci.brooklyn.cuny.edu/~sklar/teaching/f06/air/papers/smart-icra2002.pdf

Deviations from Paper

- Different task (wall following)
- Different approximation function
- Different state and action space.

Reinforcement Learning: States


```
138// converts a raw state representation to an integer representation
139int discretize_state() {
140  int root_states = floor(sqrt(NUM_STATES));
141
142  float dist_f_local = clamp(0.0, dist_f, RG_MAX_DIST);
143  int norm_f = floor((dist_f_local / RG_MAX_DIST) * (root_states - 1));
144
145  float dist_r_local = clamp(0.0, dist_r, RG_MAX_DIST);
146  int norm_r = floor((dist_r_local / RG_MAX_DIST) * (root_states - 1));
147
148  float result = norm_f + root_states * norm_r;
149
150  return result;
151}
```

Reinforcement Learning: Rewards

```
246 // determine a reward given the current state
247 // NOTE: all distances in CM (ranger values)
248 float get_reward() {
249
250   float min_f = 15.0;
251   float min_r = 15.0;
252   float max_r = 25.0;
253   float target_r = 20.0;
254
255   if (dist_f > 40) {
256     if ((old_dist_r < target_r && dist_r > old_dist_r) ||          // veered in the correct direction +1
257         (old_dist_r > target_r && dist_r < old_dist_r)) {
258
259       if (dist_r < max_r && dist_r > min_r) {                        // ended within wall range +5
260         return 10.0;
261       }
262       return 5.0;
263     }
264   }
265   if ((old_dist_r > target_r && dist_r > old_dist_r) ||          // veered in the wrong direction -5
266       (old_dist_r < target_r && dist_r < old_dist_r)) {
267     return -5;
268   }
269   if (old_dist_f > dist_f && dist_f > min_f) {                          // moved forward
270     return 1.0;
271   }
272
273   return 0.0;
274 }
```

Foreseeable Challenges

- Implementing RL
  - Training phase
- Communication between Bots

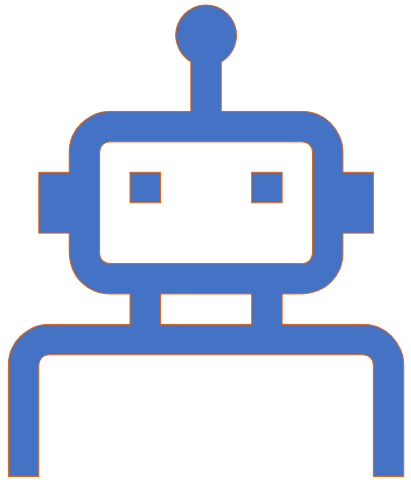

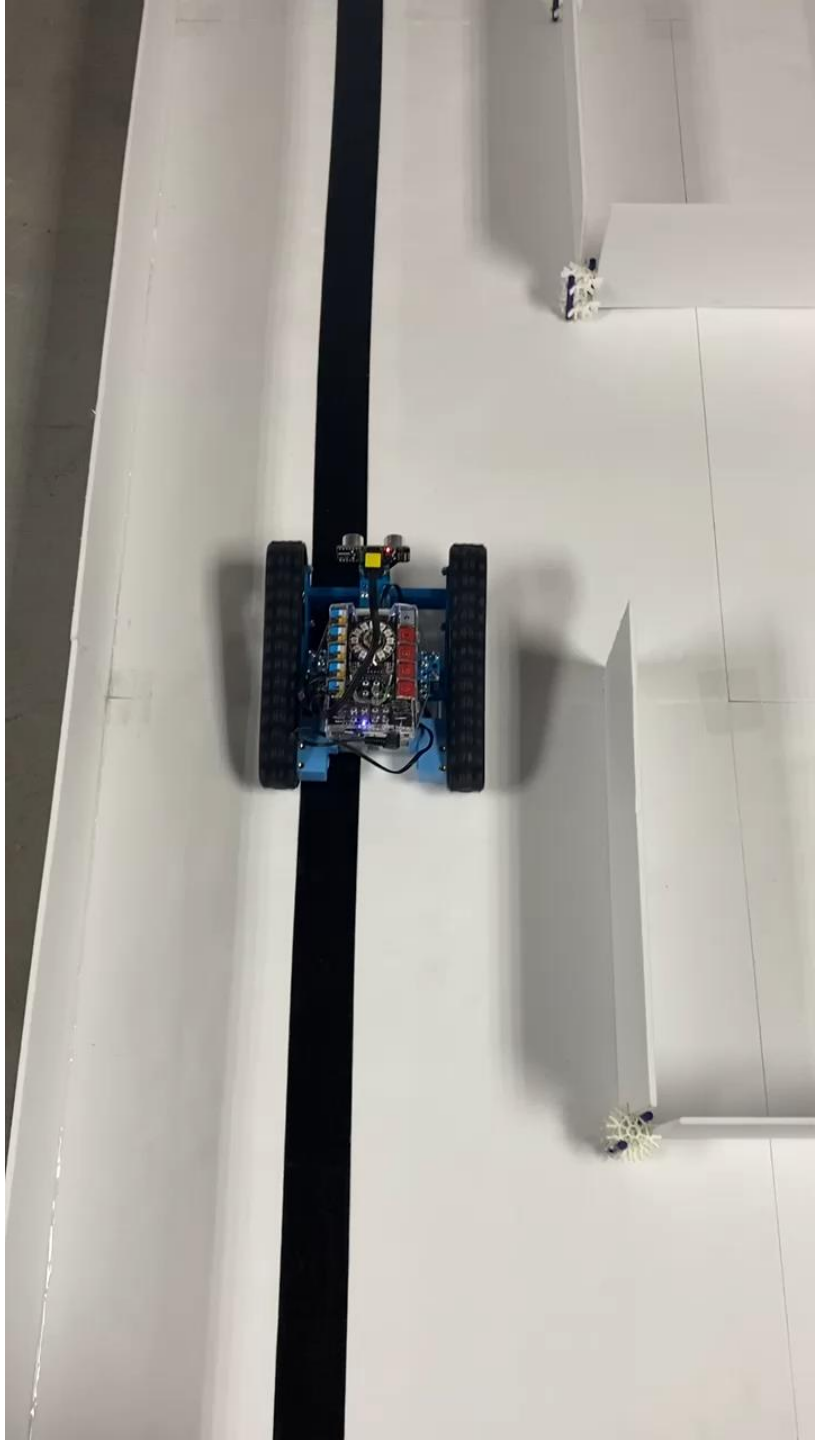

Unexpected Challenges

- Incorporating additional Sensors
- Unexpected states
- Maze Configuration

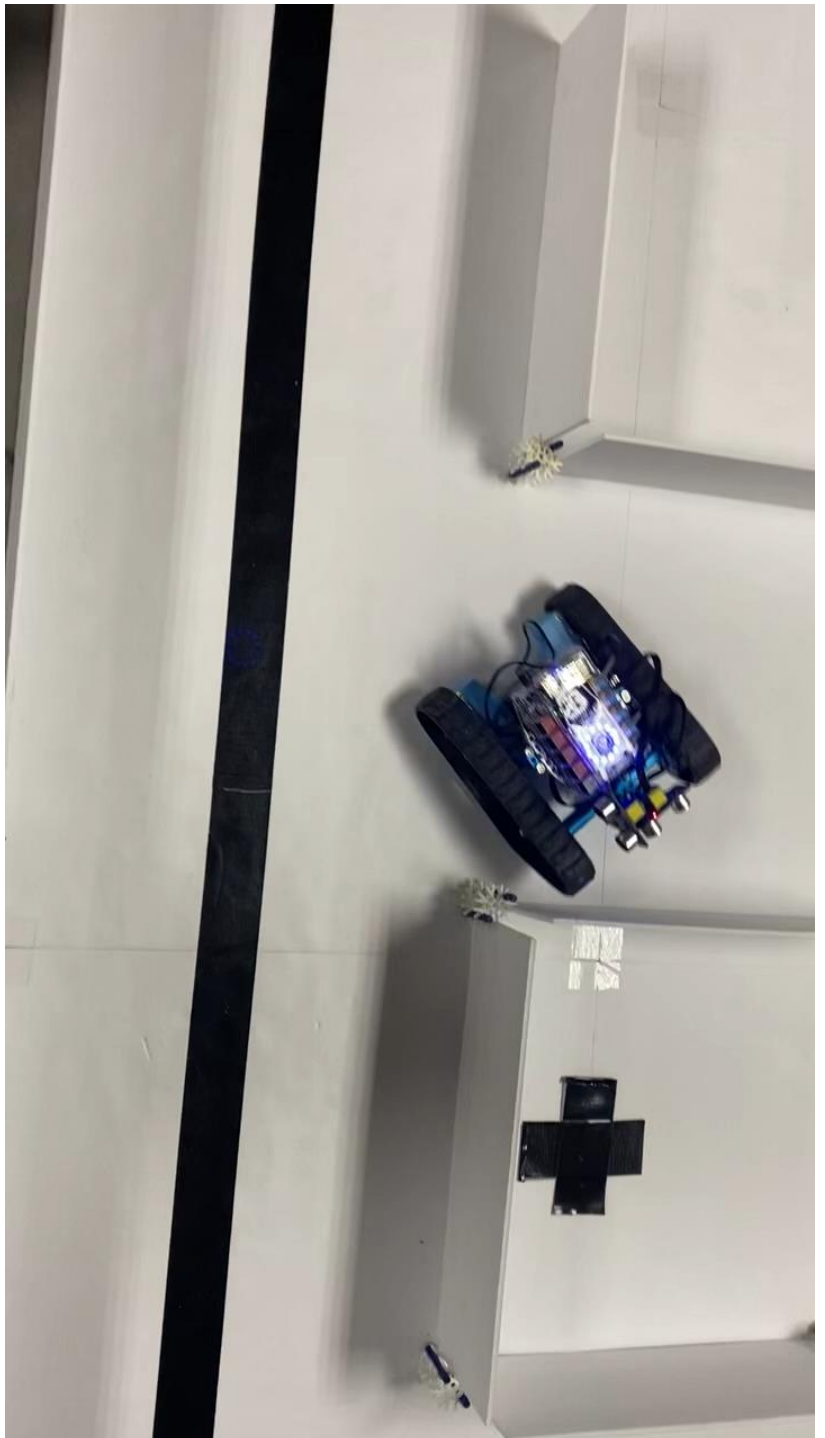

Collecting Keys

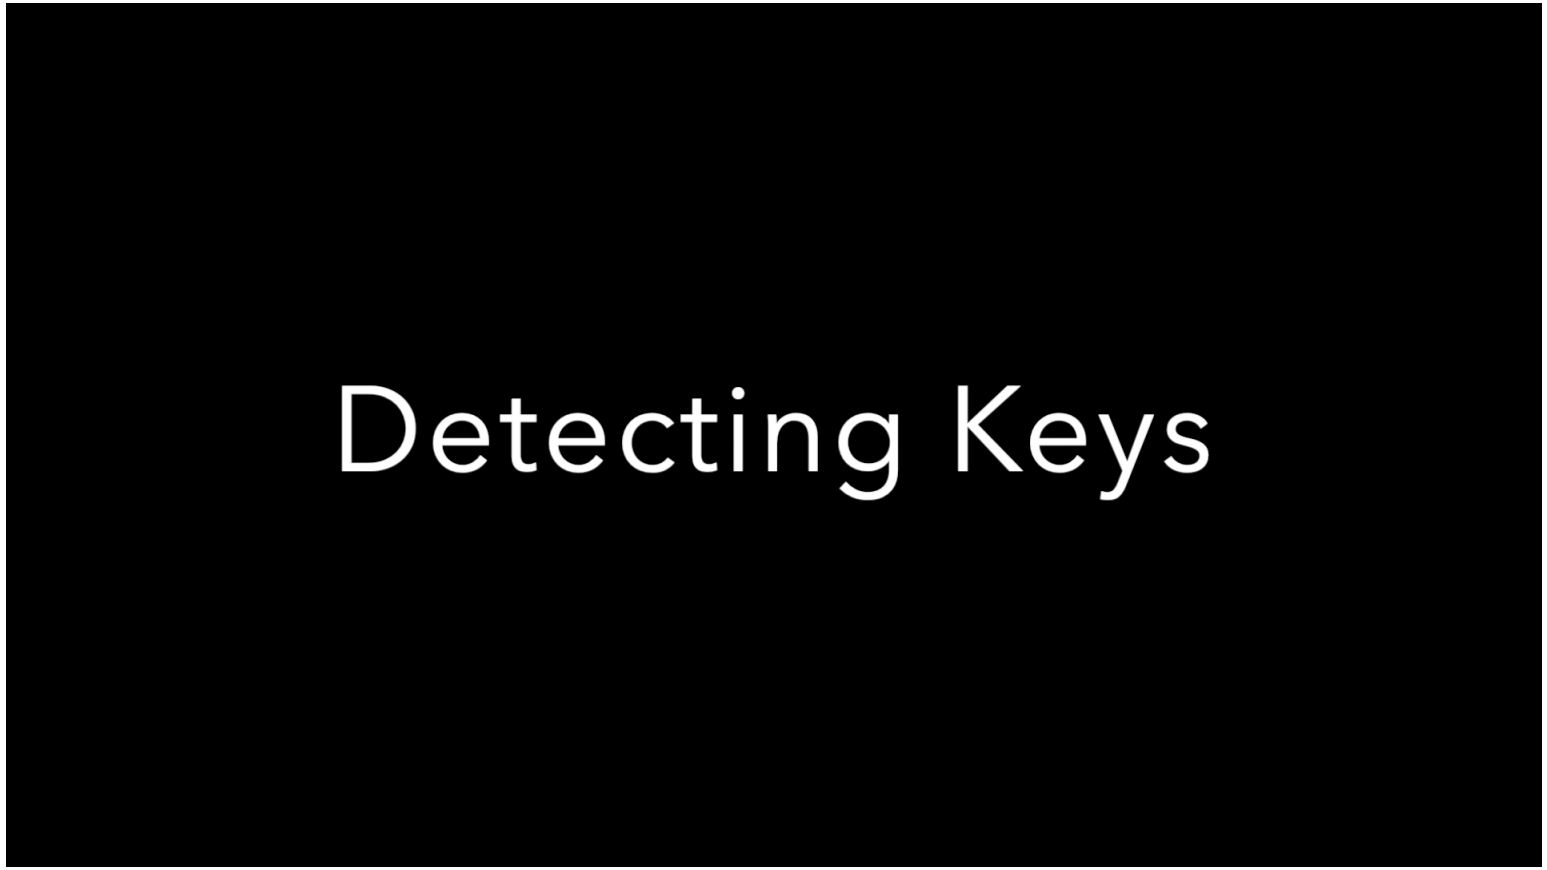

Patrol Bot

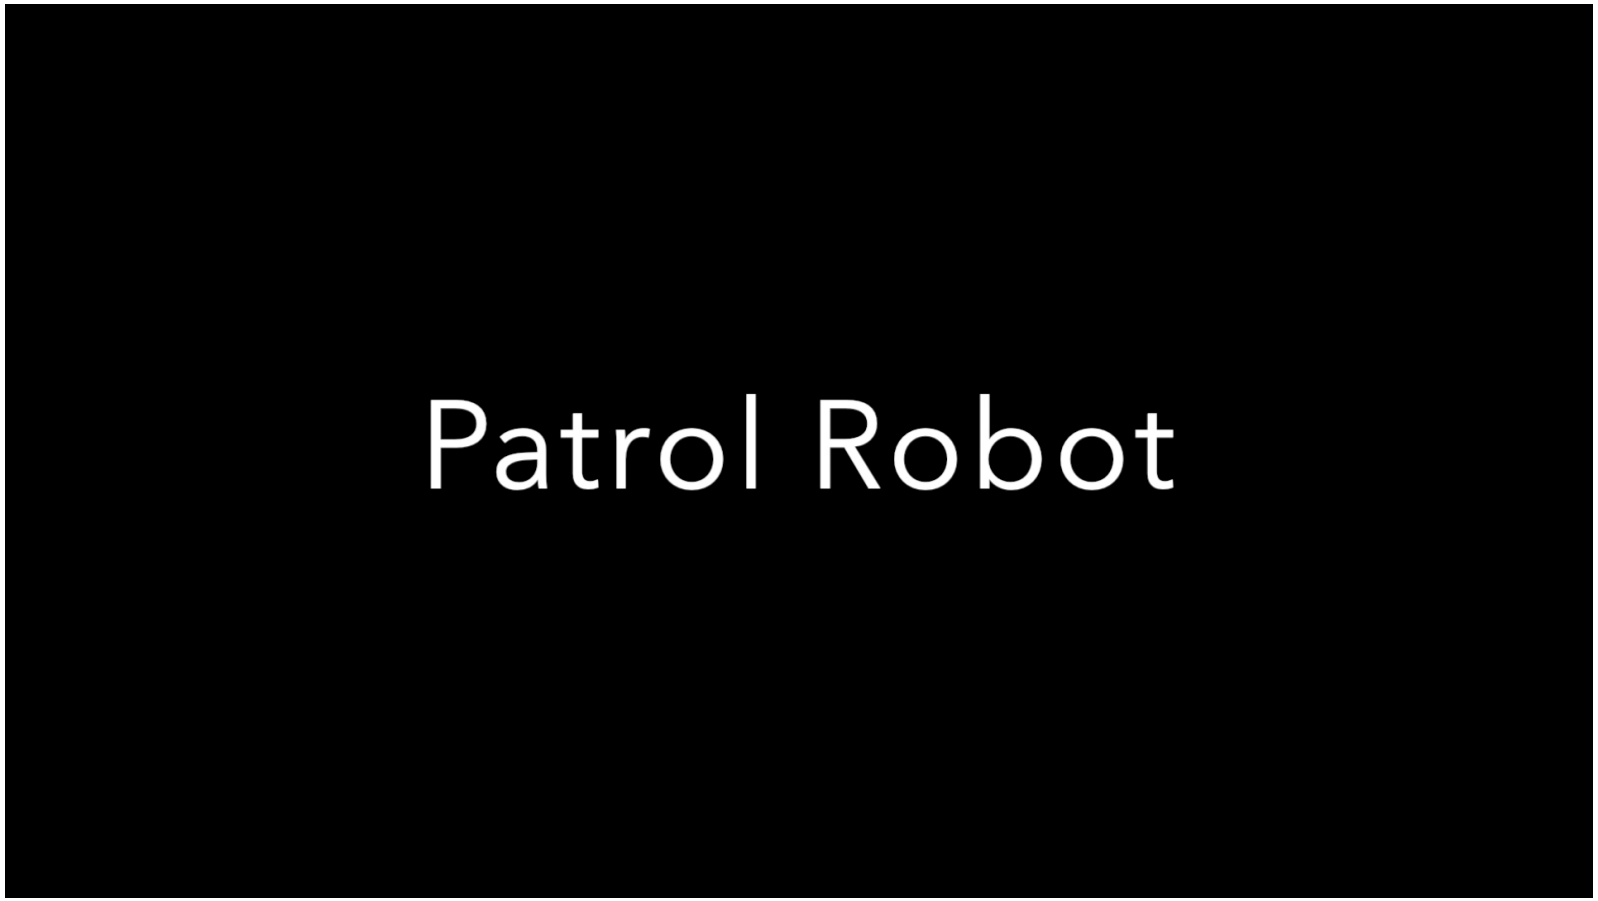

Possible Improvements

Deep learning for function approximation

Manual control for training

A physical maze more similar to gazebo

- Thicker walls
- More space in corridors